



Beyond Linear Regression Part I

Intro to Nonlinear Regression & Model Selection

송 준

고려대학교
통계학과 / 융합데이터과학 대학원

Recap: Goal of Regression Models

- 회귀 모형 (Regression Model):

$$Y = f(X) + \epsilon$$

- Goal of Regression Models:
 - 추정 (Estimation): 관계를 나타내는 함수 f 에 대한 추정
 - 예측 (Prediction): x 값이 주어졌을 때 대응되는 y 값의 예측
 - 추론 (Inference): Further investigation
 - 예측이 “얼마나” 정확한가?
 - 함수 $f()$ 가 얼마나 정확한가?
 - 예측변수가 여러 개 있을 때 모든 변수가 y 의 값에 영향을 주나?
 - 모형이 충분히 적합 됐나?

Recap: Goal of Regression Models

- 회귀 모형 (Regression Model):

$$Y = f(X) + \epsilon$$

- Goal of Regression Models:

- 추정 (Estimation): 관계를 나타내는 함수 f 에 대한 추정
- 예측 (Prediction): X 값이 x 로 주어졌을 때 Y 값의 예측
- 추론 (Inference): Further investigation of the data
 - 예측이 “얼마나” 정확한가?
 - 함수 $f()$ 가 얼마나 정확한가?
 - 예측변수가 여러 개 있을 때 모든 변수가 Y 의 값에 영향을 주나?
 - 모형이 충분히 적합 됐나?
- 예측**만 목표로 할 시: 다양한 방법론 적용 가능
- 추론**을 목표로 할 시: 관계를 나타내는 $f()$ 에 제약이 필요함
- 단순한 모형부터 시작! **f 는 선형함수.**

Recap : Function Estimation

Optimal predictor:

$$f^* = \operatorname{argmin}_f \mathbb{E}[(f(X) - Y)^2]$$

Empirical Risk Minimizer:

$$\hat{f}_n = \operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n ((f(X_i) - Y_i))^2$$

Class of predictors

Empirical mean

Recap : Linear Regression Algorithm

- **Input** : Dataset $Z = \{(x_1, y_1), \dots, (x_n, y_n)\}$
- Compute

$$\begin{aligned}\hat{\beta}(Z) &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} L(\beta; Z) \\ &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \frac{1}{n} \sum_{i=1}^n (y_i - \beta^T x_i)^2\end{aligned}$$

- **Output** : $f_{\hat{\beta}(Z)}(x) = \hat{\beta}(Z)^T x$
- Discuss algorithm for computing the minimal β later

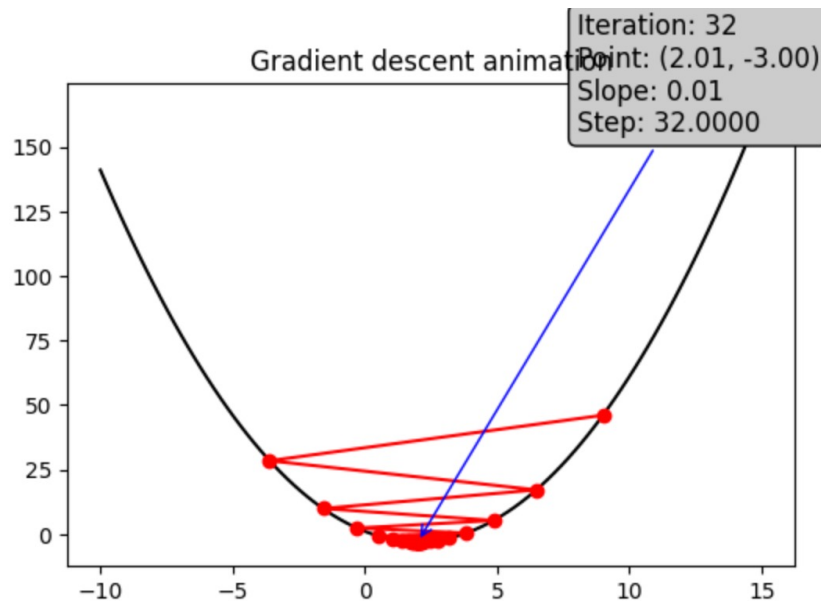
Recap : Solution to the Optimization Problem

- **Analytic Solution (Explicit form, closed form – solution): 미분=0**

$$L(\beta; x) = \beta^2 - 2x\beta + 10 = (\beta - x)^2 + 9$$

$$\arg \min_{\beta} L(\beta) = x$$

- **Numerical Solution (Optimization Algorithm)**



Recap: R^2


- SST (총 편차제곱합)

$$\sum_{i=1}^n (Y_i - \bar{Y})^2$$

결정계수 R^2 (coefficient of determination)

회귀직선의 적합도를 평가하는 방법

전체변동에서 회귀로 설명되는 부분이 차지하는 비율




$$R^2 = \frac{SSR}{SST} = 1 - \frac{SSE}{SST}$$

$$\sum_{i=1}^n (Y_i - \bar{Y})^2 = \sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2 + \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

$$\text{SST}(\text{var}) = \text{SSR} + \text{SSE} (\text{training loss})$$

Feature Mapping

Function Approximation View of ML




ML algorithm outputs a model f that best “approximates” the given data Z

Function Approximation View of ML

- *Framework for designing machine learning algorithms*
- **Two design decisions**
 - *What is the family of candidate models f ? (E.g., linear functions)*
 - *How to define “approximating”? (E.g., MSE loss)*

Aside : “True Function”

- **Input:** Dataset \mathcal{Z}
 - Presume there is an unknown function f^* that generates \mathcal{Z}
- **Goal:** Find an approximation $f_\beta \approx f^*$ in our model family $f_\beta \in F$
 - Typically, f^* not in our model family F



Function Approximation View of ML

- *Framework for designing machine learning algorithms*
- **Two design decisions**
 - *What is the family of candidate models f ? (E.g., linear functions)*
 - *How to define “approximating”? (E.g., MSE loss)*
- *How do we specialize to linear regression?*

Loss Minimization



Data Z



*Machine learning
algorithm*



Model f

Linear Regression



Data $Z = \{(x_i, y_i)\}_{i=1}^n$

$$\hat{\beta}(Z) = \underset{\beta}{\operatorname{argmin}} L(\beta; Z)$$

Model $f_{\hat{\beta}(Z)}$

L encodes $y_i \approx f_{\beta}(x_i)$

MSE loss

Model is a linear function $f_{\beta}(x) = \beta^T x$

Linear Regression

General strategy

- Model family $F = \{f_{\beta}\}_{\beta}$
- Loss function $L(\beta; Z)$

Linear regression strategy

- Linear functions $F = \{f_{\beta}(x) = \beta^T x\}$
- MSE $L(\beta; Z) = \frac{1}{n} \sum_{i=1}^n (y_i - \beta^T x_i)^2$


Linear regression algorithm

$$\hat{\beta}(Z) = \underset{\beta}{\operatorname{argmin}} L(\beta; Z)$$


Agenda

- **Function approximation view of machine learning**
 - *Modern strategy for designing machine learning algorithms*
 - **By example:** *Linear regression, a simple machine learning algorithm*
- **Bias-variance tradeoff**
 - *Fundamental challenge in machine learning*
 - **By example:** *Linear regression with feature maps*

Example: Quadratic Function



Example: Quadratic Function



Can we get a better fit?

Feature Maps

General strategy

- Model family $F = \{f_{\beta}\}_{\beta}$
- Loss function $L(\beta; Z)$

Linear regression with feature map

- Linear functions over a given **feature**

map $\phi: X \rightarrow \mathbb{R}^d$

$$F = \{f_{\beta}(x) = \beta^T \phi(x)\}$$

- MSE $L(\beta; Z) = \frac{1}{n} \sum_{i=1}^n (\mathbf{y}_i - \beta^T \phi(\mathbf{x}_i))^2$

Quadratic Feature Map


- Consider the feature map $\phi: \mathbb{R} \rightarrow \mathbb{R}^2$ given by

$$\phi(x) = \begin{bmatrix} x \\ x^2 \end{bmatrix}$$

- Then, the model family is

$$f_{\beta}(x) = \beta_1 x + \beta_2 x^2$$

Quadratic Feature Map



In our family for $\beta = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$!

Feature Maps

- *Powerful strategy for encoding prior knowledge*
- **Terminology**
 - x is the **input** and $\phi(x)$ are the **features**
 - *Often used interchangeably*

Examples of Feature Maps

- **Polynomial features**

- $f_{\beta}(x) = \beta_1 + \beta_2 x_1 + \beta_3 x_2 + \beta_4 x_1^2 + \beta_5 x_1 x_2 + \beta_6 x_2^2 + \dots$
- *Quadratic features are very common; capture “feature interactions”*
- *Can use other nonlinearities (exponential, logarithm, square root, etc.)*

- **Basis expansion approach**

- $f_{\beta}(x) = \beta_0 + \beta_1 \phi_1(x) + \dots + \beta_d \phi_d(x)$
- *Fit the data in a more general way*

- **Encoding non-real inputs**

- *E.g., x = “the food was good” and y = 4 stars*
- $\phi(x) = [1(\text{“good”} \in x) \quad 1(\text{“bad”} \in x) \quad \dots]^T$


Algorithm

- *Reduces to linear regression*
- **Step 1:** Compute $\phi_i = \phi(x_i)$ for each x_i in Z
- **Step 2:** Run linear regression with $Z' = \{(\phi_1, y_1), \dots, (\phi_n, y_n)\}$

Question


- **Why not throw in lots of features?**

- $f_{\beta}(x) = \beta_1 + \beta_2 x_1 + \beta_3 x_2 + \beta_4 x_1^2 + \beta_5 x_1 x_2 + \beta_6 x_2^2 + \dots$
- *Can fit any n points using a polynomial of degree n*



Question


- Issue: prediction with a new data
 - Given a **new** input x , predict the label $\hat{y} = f_{\beta}(x)$



The errors on new inputs is very large!

Question

- Issue: prediction with a new data
 - Given a **new** input x , predict the label $\hat{y} = f_{\beta}(x)$



Vanilla linear regression actually works better!

Model Selection Basic

Training vs. Test Data

- **Training data:** *Examples $Z = \{(x, y)\}$ used to fit our model*
- **Test data:** *New inputs x whose labels y we want to predict*
- **Goal:** *Find a model that works well on the test data (unseen data)*


Recap : Function Estimation

Ideal goal: Construct prediction rule $f^* : \mathcal{X} \rightarrow \mathcal{Y}$

$$f^* = \operatorname{argmin}_f \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$

$$\tilde{f} = \operatorname{argmin}_{f \in F} \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$

$$\widehat{f}_n = \operatorname{argmin}_{f \in F} \sum_{i=1}^n \operatorname{loss}(Y_i, f(X_i))$$



Training Loss (MSE) v.s Test Loss (MSE)

Ideal goal: Construct prediction rule $f^* : \mathcal{X} \rightarrow \mathcal{Y}$

$$f^* = \operatorname{argmin}_f \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$

$$\tilde{f} = \arg \min_{f \in F} \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$

$$\widehat{f}_n = \arg \min_{f \in F} \sum_{i=1}^n \operatorname{loss}(Y_i, f(X_i)) \quad \leftarrow \quad \text{It is obtained using the data}$$

Given the estimated function from the data, \widehat{f}_n ,

$$\textbf{Test MSE: } \mathbb{E}_{XY} \left[\operatorname{loss} \left(Y, \widehat{f}_n(X) \right) \right]$$

\leftarrow We need to minimize the Test MSE!



$$\text{Training MSE: } \frac{1}{n} \sum_{i=1}^n \operatorname{loss} \left(Y_i, \widehat{f}_n(X_i) \right)$$

Training Loss (MSE) v.s Test Loss (MSE)

- *If no test observations are available? Should we choose the method that minimizes the training MSE?*
- *No! There is no guarantee that the method with the smallest training MSE will have the smallest test MSE*

Training Loss (MSE) v.s Test Loss (MSE)


- *Training MSE (grey) decreases monotonically as the model flexibility increases and **Test MSE (red)** has U-shape*



Overfitting v.s Underfitting


- **Overfitting**

- Fit the **training data** \mathcal{Z} well
- Fit new **test data** (x, y) poorly



- **Underfitting**

- Fit the **training data** \mathcal{Z} poorly
- (Necessarily) fit new **test data** (x, y) poorly




Training Loss (MSE) vs Test Loss (MSE)

- See the [example](#)
- *There is no way to have the true Test MSE*
- *Estimate the TEST MSE!*

Training/Test Split

- **Issue:** *How to detect overfitting vs. underfitting?*
- **Solution:** Use **held-out test data** to estimate loss on new data
 - *Typically, randomly shuffle data first*



Training/Test Split Algorithm

- **Step 1:** Split Z into Z_{train} and Z_{test}

Training data Z_{train}


Test data Z_{test}

- **Step 2:** Run linear regression with Z_{train} to obtain $\hat{\beta}(Z_{train})$
- **Step 3:** Evaluate
 - **Training loss:** $L_{train} = L(\hat{\beta}(Z_{train}); Z_{train})$
 - **Estimated Test (or generalization) loss:** $L_{test} = L(\hat{\beta}(Z_{train}); Z_{test})$

Training/Test Split Algorithm


- **Overfitting**

- Fit the **training data** \mathcal{Z} well
- Fit new **test data** (x, y) poorly



- **Underfitting**


- Fit the **training data** \mathcal{Z} poorly
- (Necessarily) fit new **test data** (x, y) poorly



Training/Test Split Algorithm


- **Overfitting**

- L_{train} is small
- L_{test} is large



- **Underfitting**


- Fit the **training data** Z poorly
- (Necessarily fit new **test data** (x, y) poorly)



Training/Test Split Algorithm


- **Overfitting**

- L_{train} is small
- L_{test} is large



- **Underfitting**

- L_{train} is large
- L_{test} is large



Aside: IID Assumption

- **Underlying IID assumption**

- *Future data are drawn IID from same data distribution $P(x, y)$ as Z_{test}*
- *IID = independent and identically distributed*
- *This is a strong (but common) assumption!*

- **Time series data**

- *Particularly important failure case since data distribution may shift over time*
- **Solution:** *Split along time (e.g., data before 2024.05.18 vs. data after 2024.05.18)*

How to fix Underfitting/Overfitting?


- *Choose the right model family!*

Role of Capacity


- **Capacity** of a model family captures “complexity” of data it can fit
 - Higher capacity more likely to overfit(model family has high **variance**)
 - Lower capacity more likely to underfit(model family has high **bias**)
- For linear regression, capacity corresponds to feature dimension d
 - I.e., number of features in $\phi(x)$

Bias-Variance Tradeoff


- **Overfitting**(high **variance**)
 - *High capacity model capable of fitting complex data*
 - *Insufficient data to constrain it*



- **Underfitting**(high **bias**)
 - *Low capacity model that can only fit simple data*
 - *Sufficient data but poor fit*



Bias-Variance Tradeoff



Bias-Variance Tradeoff

- *For linear regression, increasing feature dimension d ...*
 - Tends to **increase capacity**
 - Tends to *decrease bias* but *increase variance*
- *Need to construct ϕ to balance tradeoff between bias and variance*
 - **Rule of thumb:** $n \approx d \log d$
 - *Large fraction of data science work is data cleaning + feature engineering*

Bias-Variance Tradeoff

- *Increasing number of examples n in the data...*
 - *Tends to **increase bias** and **decrease variance***
- **General strategy**
 - **High bias:** *Increase model capacity d*
 - **High variance:** *Increase data size n (i.e., gather more labeled data)*

Bias-Variance Tradeoff


Ideal goal: Construct prediction rule $f^* : \mathcal{X} \rightarrow \mathcal{Y}$

$$f^* = \operatorname{argmin}_f \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$


$$\tilde{f} = \arg \min_{f \in F} \mathbb{E}_{XY}[\operatorname{loss}(Y, f(X))]$$

$$\widehat{f}_n = \arg \min_{f \in F} \sum_{i=1}^n \operatorname{loss}(Y_i, f(X_i))$$


Overfitting?
Underfitting?




Bias-Variance Tradeoff(Overfitting)



Bias-Variance Tradeoff(Underfitting)



Bias-Variance Tradeoff(Ideal)




Cross Validation



Training/Test Split

Validation Set Approach:

Is it enough?



Validation Set Approach




- *Left: Test MSE for a single split*
- *Right: Validation method repeated 10 times, each time split is done randomly*
- *There is a lot of variability among the MSE's and this is Not good!*
- *We need more stable methods!*

Validation Set Approach

- *Advantages:*
 - *Simple and computationally less expensive.*
- *Disadvantages:*
 - *The validation MSE can be highly variable.*
 - *Only a subset of observations is used to fit the model (training data).*
 - *Statistical methods tend to perform worse when trained **on fewer observations**.*
 - *i.e., the test MSE tends to be over-estimated.*

Leave-one-out cross validation (LOOCV)



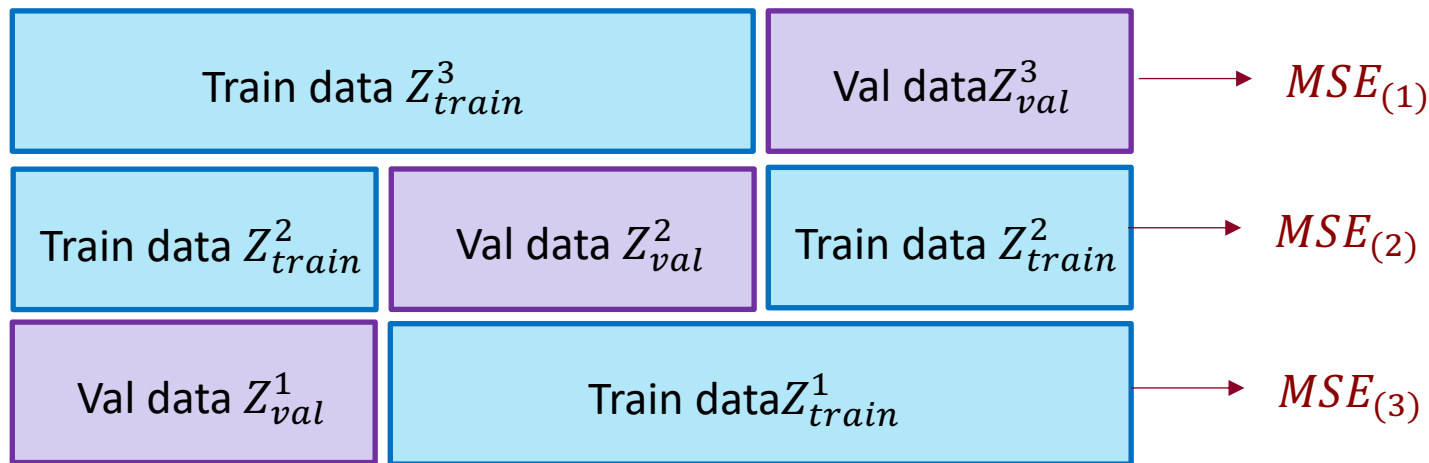
- Split the data set of size n into
 - training: $n-1$
 - validation: 1
- Fit the model using the training data
- Validate the model using the validation data: $MSE_{(i)}$ (do this n times)
- Final test MSE = $\frac{1}{n} \sum_{i=1}^n MSE_{(i)}$

Leave-one-out cross validation (LOOCV)

- *LOOCV has less bias*
 - *We repeatedly fit the statistical learning method using training data that contains $n - 1$ obs., i.e. almost all the data set is used.*
- *LOOCV is stable (no more randomness)*
- **Issue:** *LOOCV is computationally expensive*
 - *We need to fit the model n times!*

k-fold CV

- Example: K=3 fold




$$CV = \sum_{k=1}^K MSE_{(k)}$$

Choose the tuning parameter that minimizes the CV


k-fold CV

- *LOOCV v.s. 10-fold CV*

LOOCV



10-fold CV

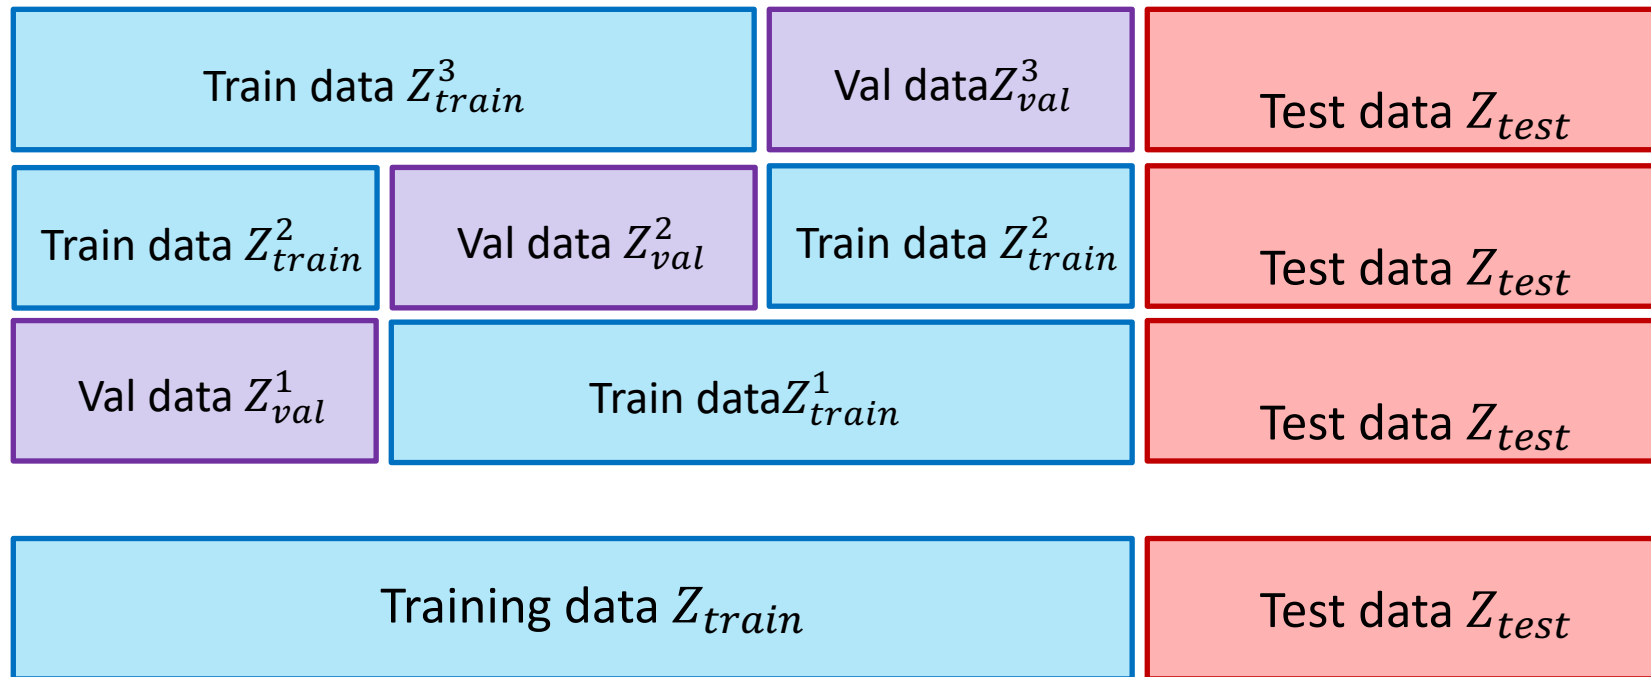


k-fold CV

- *LOOCV = n -fold CV*
- *k-fold CV*
 - *(good & bad) slightly biased / variable*
 - *(good) computationally efficient*

k-fold CV: in actual scenario

- ✓ Goal: Find a hyperparameter (model complexity, some tuning parameter)
- ✓ Split train & test. Apply k-fold CV to train!



For reporting test loss

k-Fold Cross-Validation

Goal: Find λ

- **Alternative:** k -fold cross-validation (e.g., $k = 3$)
 - Split Z into Z_{train} and Z_{test}
 - Split Z_{train} into k disjoint sets Z_{val}^s , and let $Z_{\text{train}}^s = \bigcup_{s' \neq s} Z_{\text{val}}^{s'}$
 - Use λ' that works best on average across $s \in \{1, \dots, k\}$ with Z_{train}
 - Chooses better λ' than above strategy
- **Compute vs. accuracy tradeoff**
 - As $k \rightarrow N$, the model becomes more accurate
 - But algorithm becomes more computationally expensive